

## 1 MapReduce

For each problem below, write pseudocode to complete the implementations. Tips:

- The input to each MapReduce job is given by the signature of `map()`.
- `emit(key k, value v)` outputs the key-value pair `(k, v)`.
- `for var in list` can be used to iterate through `Iterables` or you can call the `hasNext()` and `next()` functions.
- Usable data types: `int`, `float`, `String`. You may also use lists and custom data types composed of the aforementioned types.
- `intersection(list1, list2)` returns a list of the common elements of `list1`, `list2`.

- 1.1 Given a set of coins and each coin's owner in the form of a list of `CoinPairs`, compute the number of coins of each denomination that a person has.

`CoinPair`:

```
String person
String coinType
```

<pre>1 map(CoinPair pair):  map(CoinPair pair):     emit(pair, 1)</pre>	<pre>1 reduce(_____, _____):  reduce(CoinPair pair, Iterable&lt;int&gt; count):     total = 0     for num in count:         total += num     emit(pair, total)</pre>
---	--

- 1.2 Using the output of the first MapReduce, compute each person's amount of money. `valueOfCoin(String coinType)` returns a float corresponding to the dollar value of the coin.

<pre>1 map(tuple&lt;CoinPair, int&gt; output):  map(tuple&lt;CoinPair, int&gt; output):     pair, amount = output     emit(pair.person,         valueOfCoin(pair.coinType) * amount)</pre>	<pre>1 reduce(_____, _____):  reduce(String person, Iterable&lt;float&gt; values):     total = 0     for amount in values:         total += amount     emit(person, total)</pre>
--	--

## 2 Spark

**Resilient Distributed Datasets (RDD)** are the primary abstraction of a distributed collection of items

**Transforms**  $RDD \rightarrow RDD$

**map( $f$ )** Return a new transformed item formed by calling  $f$  on a source element.

**flatMap( $f$ )** Similar to map, but each input item can be mapped to 0 or more output items (so  $f$  should return a sequence rather than a single item).

**reduceByKey( $f$ )** When called on a dataset of  $(K, V)$  pairs, returns a dataset of  $(K, V)$  pairs where the values for each key are aggregated using the given reduce function  $f$ , which must be of type  $(V, V) \rightarrow V$ .

**Actions**  $RDD \rightarrow Value$

**reduce( $f$ )** Aggregate the elements of the dataset *regardless of keys* using a function  $f$ .

Call `sc.parallelize(data)` to parallelize a Python collection, `data`.

- 2.1 Given a set of coins and each coin's owner, compute the number of coins of each denomination that a person has. Then, using the output of the first result, compute each person's amount of money. Assume `valueOfCoin(coinType)` is defined and returns the dollar value of the coin.

The type of `coinPairs` is a tuple of (person, coinType) pairs.

```
1 coinData = sc.parallelize(coinPairs)

out1 = coinData.map(lambda (k1, k2): ((k1, k2), 1))
                .reduceByKey(lambda v1, v2: v1 + v2)

out2 = out1.map(lambda (k, v): (k[0], v * valueOfCoin(k[1])))
          .reduceByKey(lambda v1, v2: v1 + v2)
```

- 2.2 Given a student's name and course taken, output their name and total GPA.

CourseData:

```
int courseID
float studentGrade // a number from 0-4
```

The type of `students` is a list of (studentName, courseData) pairs.

```
1 studentsData = sc.parallelize(students)

out = studentsData.map(lambda (k, v): (k, (v.studentGrade, 1)))
                .reduceByKey(lambda v1, v2: (v1[0] + v2[0], v1[1] + v2[1]))
                .map(lambda (k, v): (k, v[0] / v[1]))
```

### 3 MapReduce/Spark Practice: Optimize the Friend Zone

- 3.1 You are given a list of tuples containing people's unique int ID and a list of the IDs of their friends. Compute the list of mutual friends between each pair of friends in a social network. You have access to the `intersection` function, which takes in two lists finds the set of elements that appear in both lists.

FriendPair:

```
int friendOne
int friendTwo
```

```
1 map(tuple<int, list<int>> info):
    map(tuple<int, list<int>> info):
        personID, friendIDs = info
        for fID in friendIDs:
            if (personID < fID):
                friendPair = (personID, fID)
            else:
                friendPair = (fID, personID)
            emit(friendPair, friendIDs)

1 reduce(_____, _____):
    reduce(FriendPair key, Iterable<list<int>> values):
        mutualFriends = intersection(
            values[0], values[1]
        )
        emit(key, mutualFriends)
```

- 3.2 Solve the problem above using Spark.

The type of `persons` is a list of `(personID, list(friendID))` pairs.

```
1 def genFriendPairAndValue(pair):
2     pID, fIDs = pair
3     return [(pID, fID), fIDs] if pID < fID else [(fID, pID), fIDs] for fID in fIDs
4
5 def intersection(l1, l2):
6     return [x for x in l1 if x in l2]
7
8 personsData = sc.parallelize(persons)
```

```
out = personsData.flatMap(genFriendPairAndValue).reduceByKey(intersection)
```

## 4 Warehouse-Scale Computing

Sources speculate Google has over 1 million servers. Assume each of the 1 million servers draw an average of 200W, the PUE is 1.5, and that Google pays an average of 6 cents per kilowatt-hour for datacenter electricity.

- 4.1 Estimate Google's annual power bill for its datacenters.

$1.5 \cdot 10^6 \text{ servers} \cdot 0.2\text{kW/server} \cdot \$0.06/\text{kW-hr} \cdot 8760 \text{ hrs/yr} \approx \$157.68 \text{ M/year}$

- 4.2 Google reduced the PUE of a 50,000-machine datacenter from 1.5 to 1.25 without decreasing the power supplied to the servers. What's the cost savings per year?

$$\text{PUE} = \frac{\text{Total building power}}{\text{IT equipment power}} \implies \text{Savings} \propto (\text{PUE}_{\text{old}} - \text{PUE}_{\text{new}}) \cdot \text{IT equipment power}$$

$$(1.5 - 1.25) \cdot 50000 \text{ servers} \cdot 0.2 \text{ kW/server} \cdot \$0.06/\text{kW-hr} \cdot 8760 \text{ hrs/yr} \approx \$1.314 \text{ M/year}$$